

Data technologies, formats, and structures

Chris Paciorek

2022-08-23

Table of contents

1. Data storage and file formats on a computer	2
Text and binary files	2
Common file types	3
2. Reading data from text files into R	4
Core R functions	4
Connections	7
File paths	9
The <i>readr</i> package	10
Reading data quickly	11
3. Output from R	11
Writing output to files	11
Formatting output	11
4. Webscraping and working with HTML, XML, and JSON	13
Reading HTML	13
XML	16
JSON	20
Webscraping and web APIs	23
Webscraping ethics and best practices	23
What is HTTP?	24
APIs: REST- and SOAP-based web services	24
HTTP requests by deconstructing an (undocumented) API	26
More details on HTTP requests	28
Packaged access to an API	29
Accessing dynamic pages	30
5. File and string encodings	31
6. Data structures	35
Standard data structures in R and Python	35
Other kinds of data structures	35

PDF

References (see [syllabus](#) for links):

- Adler
- Nolan and Temple Lang, XML and Web Technologies for Data Sciences with R.
- Chambers
- [R intro manual](#) on CRAN (R-intro).
- Venables and Ripley, Modern Applied Statistics with S
- Murrell, Introduction to Data Technologies.
- [R Data Import/Export manual](#) on CRAN (R-data).
- SCF tutorial: [Working with large datasets in SQL, R, and Python](#)

(Optional) Videos

There are four videos from 2020 in the bCourses Media Gallery that you can use for reference if you want to:

1. Text files and ASCII
2. Encodings and UTF-8
3. HTML
4. XML and JSON

1. Data storage and file formats on a computer

We're going to start early in the data analysis pipeline: getting data, reading data in, writing data out to disk, and webscraping. We'll focus on doing these manipulations in R, but the concepts and tools involved are common to other languages, so familiarity with these in R should allow you to pick up other tools easily. The main downside to working with datasets in R (true for Python and most other languages as well) is that the entire dataset resides in memory, so R is not so good for dealing with very large datasets. More on alternatives in a later unit. R (and similar languages) has the capability to read in a wide variety of file formats.

Text and binary files

In general, files can be divided into text files and binary files. In both cases, information is stored as a series of bits. Recall that a bit is a single value in base 2 (i.e., a 0 or a 1), while a byte is 8 bits.

A **text file** is one in which the bits in the file encode individual characters. Note that the characters can include the digit characters 0-9, so one can include numbers in a text file by writing down the digits needed for the number of interest. Examples of text file formats include CSV, XML, HTML, and JSON.

Text files may be simple ASCII files (i.e., files encoded using ASCII) or in other encodings such as UTF-8, both covered in Section 5. [ASCII](#) files have 8 bits (1 byte) per character and can represent 128 characters (the 52 lower and upper case letters in English, 10 digits, punctuation and a few other things – basically what you see on a standard US keyboard). UTF-8 files have between 1 and 4 bytes per character.

A **binary file** is one in which the bits in the file encode the information in a custom format and not simply individual characters. Binary formats are not (easily) human readable but can be more space-efficient and faster to work with (because it can allow random access into the data rather than requiring sequential reading). The meaning of the bytes in such files depends on the specific binary format being used and a program that uses the file needs to know how the format represents information. Examples of binary files include netCDF files, R data (e.g., .Rda) files, Python pickle files, and compiled code files.

Numbers in binary files are usually stored as 8 bytes per number. We'll discuss this much more in Unit 8.

Common file types

Here are some of the common file types, some of which are text formats and some of which are binary formats.

1. 'Flat' text files: data are often provided as simple text files. Often one has one record or observation per row and each column or field is a different variable or type of information about the record. Such files can either have a fixed number of characters in each field (*fixed width format*) or a special character (a *delimiter*) that separates the fields in each row. Common delimiters are tabs, commas, one or more spaces, and the pipe (`|`). Common file extensions are *.txt* and *.csv*. Metadata (information about the data) are often stored in a separate file. CSV files are quite common, but if you have files where the data contain commas, other delimiters can be good. Text can be put in quotes in CSV files, and this can allow use of commas within the data. This is difficult to deal with from the command line, but *read.table()* in R handles this situation.
 - One occasionally tricky difficulty is as follows. If you have a text file created in Windows, the line endings are coded differently than in UNIX. Windows uses a newline (the ASCII character `\n`) and a carriage return (the ASCII character `\r`) whereas UNIX uses only a newline in UNIX). There are UNIX utilities (*fromdos* in Ubuntu, including the SCF Linux machines and *dos2unix* in other Linux distributions) that can do the necessary conversion. If you see `^M` at the end of the lines in a file, that's the tool you need. Alternatively, if you open a UNIX file in Windows, it may treat all the lines as a single line. You can fix this with *todos* or *unix2dos*.
2. In some contexts, such as textual data and bioinformatics data, the data may be in a text file with one piece of information per row, but without meaningful columns/fields.
3. In scientific contexts, netCDF (*.nc*) (and the related HDF5) are popular format for gridded data that allows for highly-efficient storage and contains the metadata within the file. The basic structure of a netCDF file is that each variable is an array with multiple dimensions (e.g., latitude, longitude, and time), and one can also extract the values of and metadata about each dimension. The *ncdf4* package in R nicely handles working with netCDF files.
4. Data may also be in text files in formats designed for data interchange between various languages, in particular XML or JSON. These formats are "self-describing"; namely the metadata is part of the file. The *XML2*, *rvest*, and *jsonlite* packages are useful for reading and writing from these formats. More in Section 4.

5. You may be scraping information on the web, so dealing with text files in various formats, including HTML. The *XML2* and *rvest* packages are also useful for reading HTML.
6. Data may already be in a database or in the data storage format of another statistical package (*Stata*, *SAS*, *SPSS*, etc.). The *foreign* package in R has excellent capabilities for importing *Stata* (*read.dta()*), *SPSS* (*read.spss()*), and *SAS* (*read.ssd()*) and, for XPORT files, *read.xport()*, among others.
7. For Excel, there are capabilities to read an Excel file (see the *readxl* and *XLConnect* package among others), but you can also just go into Excel and export as a CSV file or the like and then read that into R. In general, it's best not to pass around data files as Excel or other spreadsheet format files because (1) Excel is proprietary, so someone may not have Excel and the format is subject to change, (2) Excel imposes limits on the number of rows, (3) one can easily manipulate text files such as CSV using UNIX tools, but this is not possible with an Excel file, (4) Excel files often have more than one sheet, graphs, macros, etc., so they're not a data storage format per se.
8. R can easily interact with databases (SQLite, PostgreSQL, MySQL, Oracle, etc.), querying the database using SQL and returning results to R. More in the big data unit and in the large datasets tutorial mentioned above.

2. Reading data from text files into R

Core R functions

read.table() is probably the most commonly-used function for reading in data. It reads in delimited files (*read.csv()* and *read.delim()* are special cases of *read.table()*). The key arguments are the delimiter (the *sep* argument) and whether the file contains a header, a line with the variable names. We can use *read.fwf()* to read from a fixed width text file into a data frame.

The most difficult part of reading in such files can be dealing with how R determines the classes of the fields that are read in. There are a number of arguments to *read.table()* and *read.fwf()* that allow the user to control the classes. One difficulty in older versions of R was that character fields were read in as factors.

Let's work through a couple examples. Before we do that, let's look at the arguments to *read.table()*. Note that *sep=' '* separates on any amount of white space.

```
dat <- read.table(file.path('..', 'data', 'RTADataSub.csv'),
                  sep = ',', header = TRUE)
sapply(dat, class)[1:10] # What are the classes of the columns?
```

X2010.08.02.18.55	X2336	X549	X2086
"character"	"character"	"character"	"character"
X666	X481	X298	X1624
"character"	"character"	"character"	"character"
X1732	X593		
"character"	"character"		

```
## whoops, there is an 'x', presumably indicating missingness:
unique(dat[, 2])
```

```
[1] "2124" "1830" "1833" "1600" "1578" "1187" "1005" "918" "865" "871"
[11] "860" "883" "897" "898" "893" "913" "870" "962" "880" "875"
[21] "884" "894" "836" "848" "885" "851" "900" "861" "866" "867"
[31] "829" "853" "920" "877" "908" "855" "845" "859" "856" "825"
[41] "828" "854" "847" "840" "873" "822" "818" "838" "815" "813"
[51] "816" "849" "802" "805" "792" "823" "808" "798" "800" "842"
[61] "809" "807" "826" "810" "801" "794" "771" "796" "790" "787"
[71] "775" "751" "783" "811" "768" "779" "795" "770" "821" "830"
[81] "767" "772" "791" "781" "773" "777" "814" "778" "782" "837"
[91] "759" "846" "797" "835" "832" "793" "803" "834" "785" "831"
[101] "820" "812" "824" "728" "760" "762" "753" "758" "764" "741"
[111] "709" "735" "749" "752" "761" "750" "776" "766" "789" "763"
[121] "864" "858" "869" "886" "844" "863" "916" "890" "872" "907"
[131] "926" "935" "933" "906" "905" "912" "972" "996" "1009" "961"
[141] "952" "981" "917" "1011" "1071" "1920" "3245" "3805" "3926" "3284"
[151] "2700" "2347" "2078" "2935" "3040" "1860" "1437" "1512" "1720" "1493"
[161] "1026" "928" "874" "833" "850" "" "x"
```

```
## let's treat 'x' as a missing value indicator
dat2 <- read.table(file.path('..', 'data', 'RTADataSub.csv'),
  sep = ',', header = TRUE,
  na.strings = c("NA", "x"))
unique(dat2[, 2])
```

```
[1] 2124 1830 1833 1600 1578 1187 1005 918 865 871 860 883 897 898 893
[16] 913 870 962 880 875 884 894 836 848 885 851 900 861 866 867
[31] 829 853 920 877 908 855 845 859 856 825 828 854 847 840 873
[46] 822 818 838 815 813 816 849 802 805 792 823 808 798 800 842
[61] 809 807 826 810 801 794 771 796 790 787 775 751 783 811 768
[76] 779 795 770 821 830 767 772 791 781 773 777 814 778 782 837
[91] 759 846 797 835 832 793 803 834 785 831 820 812 824 728 760
[106] 762 753 758 764 741 709 735 749 752 761 750 776 766 789 763
[121] 864 858 869 886 844 863 916 890 872 907 926 935 933 906 905
[136] 912 972 996 1009 961 952 981 917 1011 1071 1920 3245 3805 3926 3284
[151] 2700 2347 2078 2935 3040 1860 1437 1512 1720 1493 1026 928 874 833 850
[166] NA
```

```
## Let's check that the empty strings from 'dat' are now NAs in 'dat2'
which(dat[, 2] == "")[1:10]
```

```
[1] 312 313 314 315 317 318 319 320 322 323
```

```
dat2[which(dat[, 2] == "")[1], ] # pull out a line with a missing string
```

```

      X2010.08.02.18.55 X2336 X549 X2086 X666 X481 X298 X1624 X1732 X593 X222
312  2010-08-03 10:31    NA  NA    NA  NA  NA  NA    NA  NA    NA  NA    NA
      X911 X261 X1730 X211 X365 X216 X438 X596 X206 X204 X270 X176 X1159 X1137
312  NA  NA    NA  NA  NA  NA  NA  NA  NA  NA  NA  NA    NA  NA
      X135 X2036 X138 X1038 X201 X610 X627 X195 X976 X151 X1830 X421 X1087 X1157
312  NA  NA    NA  NA  NA  NA  NA  NA  NA  NA  NA  NA    NA  NA
      X181 X267 X193 X391 X208 X614 X546 X186 X1391 X217 X230 X625 X376 X164 X329
312  NA  NA    NA  NA  NA  NA  NA  NA  NA  NA  NA  NA    NA  NA
      X1043 X497 X440 X197 X287 X837 X226 X973
312  NA  NA    NA  NA  NA  NA  NA  NA  NA

```

Using `colClasses` is a good way to control how data are read in.

```

sequ <- read.table(file.path('.', 'data', 'hivSequ.csv'),
  sep = ',', header = TRUE,
  colClasses = c('integer','integer','character',
    'character','numeric','integer'))
## let's make sure the coercion worked - sometimes R is obstinant
sapply(sequ, class)

```

```

PatientID      Resp      PR.Seq      RT.Seq      VL.t0      CD4.t0
"integer"      "integer" "character" "character"  "numeric"  "integer"

```

```
## that made use of the fact that a data frame is a list
```

Note that you can avoid reading in one or more columns by specifying *NULL* as the column class for those columns to be omitted. Also, specifying the *colClasses* argument explicitly should make for faster file reading. Finally, setting `stringsAsFactors=FALSE` is standard practice and is the default in R as of version 4.0. (*readr::read_csv()* has always set `stringsAsFactors=FALSE`).

If possible, it's a good idea to look through the input file in the shell or in an editor before reading into R to catch such issues in advance. Using the UNIX command *less* on *RTADataSub.csv* would have revealed these various issues, but note that *RTADataSub.csv* is a 1000-line subset of a much larger file of data available from the kaggle.com website. So more sophisticated use of UNIX utilities (as we will see in Unit 3) is often useful before trying to read something into a program.

The basic function *scan()* simply reads everything in, ignoring lines, which works well and very quickly if you are reading in a numeric vector or matrix. *scan()* is also useful if your file is free format - i.e., if it's not one line per observation, but just all the data one value after another; in this case you can use *scan()* to read it in and then format the resulting character or numeric vector as a matrix with as many columns as fields in the dataset. Remember that the default is to fill the matrix by column.

If the file is not nicely arranged by field (e.g., if it has ragged lines), we'll need to do some more work. *readLines()* will read in each line into a separate character vector, after which we can process the lines

using text manipulation. Here's an example from some US meteorological data where I know from metadata (not provided here) that the 4-11th values are an identifier, the 17-20th are the year, the 22-23rd the month, etc.

```
dat <- readLines(file.path('.', 'data', 'precip.txt'))
id <- as.factor(substring(dat, 4, 11) )
year <- substring(dat, 18, 21)
year[1:5]
```

```
[1] "2010" "2010" "2010" "2010" "2010"
```

```
class(year)
```

```
[1] "character"
```

```
year <- as.integer(substring(dat, 18, 21))
month <- as.integer(substring(dat, 22, 23))
nvalues <- as.integer(substring(dat, 28, 30))
```

Actually, that file, *precip.txt*, is in a fixed-width format (i.e., every element in a given column has the exact same number of characters), so reading in using *read.fwf()* would be a good strategy.

Connections

R allows you to read in not just from a file but from a more general construct called a *connection*. This can include reading in text from the output of running a shell command and from unzipping a file on the fly.

Here are some examples of connections:

```
dat <- readLines(pipe("ls -al"))
dat <- read.table(pipe("unzip dat.zip"))
dat <- read.csv(gzfile("dat.csv.gz"))
dat <- readLines("http://www.stat.berkeley.edu/~paciorek/index.html")
```

In some cases, you might need to create the connection using *url()* or using the *curl()* function from the *curl* package. Though for the example here, simply passing the URL to *readLines()* does work. (In general, *curl::curl()* provides some nice features for reading off the internet.)

```
wikip1 <- readLines("https://wikipedia.org")
wikip2 <- readLines(url("https://wikipedia.org"))
library(curl)
```

Using libcurl 7.68.0 with GnuTLS/3.6.13

```
wikip3 <- readLines(curl("https://wikipedia.org"))
```

If a file is large, we may want to read it in chunks (of lines), do some computations to reduce the size of things, and iterate. This is referred to as online processing. `read.table()`, `read.fwf()` and `readLines()` all have the arguments that let you read in a fixed number of lines. To read-on-the-fly in blocks, we need to first establish the connection and then read from it sequentially. (If you don't, you'll read from the start of the file every time you read from the file.)

```
con <- file(file.path(".", "data", "precip.txt"), "r")
## "r" for 'read' - you can also open files for writing with "w"
## (or "a" for appending)
class(con)
blockSize <- 1000 # obviously this would be large in any real application
nLines <- 300000
for(i in 1:ceiling(nLines / blockSize)){
  lines <- readLines(con, n = blockSize)
  # manipulate the lines and store the key stuff
}
close(con)
```

Here's an example of using `curl()` to do this for a file on the web.

```
URL <- "https://www.stat.berkeley.edu/share/paciorek/2008.csv.gz"
con <- gzcon(curl(URL, open = "r"))
## url() in place of curl() works too
for(i in 1:4) { ## Read in first four chunks as an example
  print(i)
  print(system.time(tmp <- readLines(con, n = 100000)))
  print(tmp[1])
}
```

```
[1] 1
   user system elapsed
0.449  0.000  0.450
[1] "Year,Month,DayofMonth,DayOfWeek,DepTime,CRSDepTime,ArrTime,CRSArrTime,UniqueCarrier,FlightNum,Ta
[1] 2
   user system elapsed
0.436  0.000  0.437
[1] "2008,1,29,2,1938,1935,2308,2257,XE,7676,N11176,150,142,104,11,3,SLC,OKC,866,5,41,0,,0,NA,NA,NA,N
[1] 3
   user system elapsed
0.439  0.000  0.439
[1] "2008,1,20,7,1540,1525,1651,1637,00,5703,N227SW,71,72,58,14,15,SBA,SJC,234,5,8,0,,0,NA,NA,NA,NA,N
```



```
[1] 4
      user  system elapsed
0.424   0.000   0.424
[1] "2008,1,2,3,1313,1250,1443,1425,WN,440,N461WN,150,155,138,18,23,MC0,STL,880,3,9,0,,0,2,0,0,0,16"

close(con)
```

More details on sequential (on-line) processing of large files can be found in the tutorial on large datasets mentioned in the reference list above.

One cool trick that can come in handy is to create a *text connection*. This lets you ‘read’ from an R character vector as if it were a text file and could be handy for processing text. For example, you could then use *read.fwf()* applied to *con*.

```
dat <- readLines('../data/precip.txt')
con <- textConnection(dat[1], "r")
read.fwf(con, c(3,8,4,2,4,2))
```

```
      V1      V2   V3 V4   V5 V6
1 DLY 1000807 PRCP HI 2010  2
```

We can create connections for writing output too. Just make sure to open the connection first.

File paths

A few notes on file paths, related to ideas of reproducibility.

1. In general, you don’t want to hard-code absolute paths into your code files because those absolute paths won’t be available on the machines of anyone you share the code with. Instead, use paths relative to the directory the code file is in, or relative to a baseline directory for the project, e.g.:

```
dat <- read.csv('../data/cpds.csv')
```

2. Be careful with the directory separator in Windows files: you can either do `C:\\mydir\\file.txt` or `C:/mydir/file.txt`, but not `C:\mydir\file.txt`, and note the next comment about avoiding use of ‘\’ for portability.
3. Using UNIX style directory separators will work in Windows, Mac or Linux, but using Windows style separators is not portable across operating systems.

```
## good: will work on Windows
dat <- read.csv('../data/cpds.csv')
## bad: won't work on Mac or Linux
dat <- read.csv('..\\data\\cpds.csv')
```

4. Even better, use `file.path()` so that paths are constructed specifically for the operating system the user is using:

```
## good: operating-system independent
dat <- read.csv(file.path('..', 'data', 'cpds.csv'))
```

The *readr* package

readr is intended to deal with some of the shortcomings of the base R functions, such as leaving column names unmodified, and recognizing dates/times. It reads data in much more quickly than the base R equivalents. See [this blog post](#). Some of the *readr* functions that are analogs to the comparably-named base R functions are `read_csv()`, `read_fwf()`, `read_lines()`, and `read_table()`.

Let's try out `read_csv()` on the airline dataset used in the R bootcamp.

```
library(readr)
```

Attaching package: 'readr'

The following object is masked from 'package:curl':

```
parse_date
```

```
## I'm violating the rule about absolute paths here!!
## (airline.csv is big enough that I don't want to put it in the
##    course repository)
dir <- "~/staff/workshops/r-bootcamp-fall-2020/data"
system.time(dat <- read.csv(file.path(dir, 'airline.csv'), stringsAsFactors = FALSE))
```

```
user  system elapsed
3.537   0.145   3.685
```

```
system.time(dat2 <- read_csv(file.path(dir, 'airline.csv')))
```

Rows: 539895 Columns: 29

-- Column specification -----

Delimiter: ","

chr (5): UniqueCarrier, TailNum, Origin, Dest, CancellationCode

dbl (24): Year, Month, DayOfMonth, DayOfWeek, DepTime, CRSDepTime, ArrTime, ...

i Use ``spec()`` to retrieve the full column specification for this data.

i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

```
user  system elapsed
```

3.794 0.129 1.194

Reading data quickly

In addition to the tips above, there are a number of packages that allow one to read large data files quickly, in particular *data.table*, *arrow*, and *fst*. In general, these provide the ability to load datasets into R without having them in memory, but rather stored in clever ways on disk that allow for fast access. Metadata is stored in R. More on this in the unit on big data and in the tutorial on large datasets mentioned in the reference list above.

3. Output from R

Writing output to files

Functions for text output are generally analogous to those for input. *write.table()*, *write.csv()*, and *writeLines()* are analogs of *read.table()*, *read.csv()*, and *readLines()*. *write_csv()* is the *readr* version of *write.csv*. *write()* can be used to write a matrix to a file, specifying the number of columns desired. *cat()* can be used when you want fine control of the format of what is written out and allows for outputting to a connection (e.g., a file).

toJSON() in the *jsonlite* package will output R objects as JSON. One use of JSON as output from R would be to *serialize* the information in an R object such that it could be read into another program.

And of course you can always save to an R data file (a binary file format) using *save.image()* (to save all the objects in the workspace or *save()* to save only some objects. Happily this is platform-independent so can be used to transfer R objects between different OS.

Formatting output

cat() is a good choice for printing a message to the screen, often better than *print()*, which is an object-oriented method. You generally won't have control over how the output of a *print()* statement is actually printed.

```
val <- 1.5
cat('My value is ', val, '.\n', sep = '')
```

My value is 1.5.

```
print(paste('My value is ', val, '.', sep = ''))
```

[1] "My value is 1.5."

We can do more to control formatting with *cat()*:

```
## input
x <- 7
```

```
n <- 5
## display powers
cat("Powers of", x, "\n")
```

Powers of 7

```
cat("exponent  result\n\n")
```

exponent result

```
result <- 1
for (i in 1:n) {
  result <- result * x
  cat(format(i, width = 8), format(result, width = 10),
      "\n", sep = "")
}
```

```
1      7
2     49
3    343
4   2401
5  16807
```

One thing to be aware of when writing out numerical data is how many digits are included. For example, the default with `write` and `cat` is the number of digits that R displays to the screen, controlled by `options()$digits`. But note that `options()$digits` seems to have some variability in behavior across operating systems. If you want finer control, use `sprintf`. For example, here we print out temperatures as reals (“f”=floating point) with four decimal places and nine total character positions, followed by a C for Celsius:

```
temps <- c(12.5, 37.234324, 1342434324.79997234, 2.3456e-6, 1e10)
sprintf("%9.4f C", temps)
```

```
[1] " 12.5000 C"      " 37.2343 C"      "1342434324.8000 C"
[4] "  0.0000 C"      "10000000000.0000 C"
```

```
city <- "Boston"
sprintf("The temperature in %s was %.4f C.", city, temps[1])
```

```
[1] "The temperature in Boston was 12.5000 C."
```

```
sprintf("The temperature in %s was %9.4f C.", city, temps[1])
```

```
[1] "The temperature in Boston was 12.5000 C."
```

To change the number of digits printed to the screen, do `options(digits = 5)` or specify as an argument to `print` or use `sprintf`.

4. Webscraping and working with HTML, XML, and JSON

The book *XML and Web Technologies for Data Sciences with R* by Deb Nolan (UCB Stats faculty) and Duncan Temple Lang (UCB Stats PhD alumnus and UC Davis Stats faculty) provides extensive information about getting and processing data off of the web, including interacting with web services such as REST and SOAP and programmatically handling authentication.

Here are some UNIX command-line tools to help in webscraping and working with files in formats such as JSON, XML, and HTML: <http://jeroenjanssens.com/2013/09/19/seven-command-line-tools-for-data-science.html>.

We'll cover a few basic examples in this section, but HTML and XML formatting and navigating the structure of such pages in great detail is beyond the scope of what we can cover. The key thing is to see the main concepts and know that the tools exist so that you can learn how to use them if faced with such formats.

Reading HTML

HTML (Hypertext Markup Language) is the standard markup language used for displaying content in a web browser. In simple webpages (ignoring the more complicated pages that involve Javascript), what you see in your browser is simply a *rendering* (by the browser) of a text file containing HTML.

However, instead of rendering the HTML in a browser, we might want to use code to extract information from the HTML.

Let's see a brief example of reading in HTML tables.

Note that before doing any coding, it can be helpful to look at the raw HTML source code for a given page. We can explore the underlying HTML source in advance of writing our code by looking at the page source directly in the browser (e.g., in Firefox under the 3-lines “open menu” symbol, see **Web Developer (or More Tools) -> Page Source** and in Chrome **View -> Developer -> View Source**), or by downloading the webpage and looking at it in an editor, although in some cases (such as the `nytimes.com` case), what we might see is a lot of JavaScript.

One lesson here is not to write a lot of your own code to do something that someone else has probably already written a package for. We'll use the *rvest* package.

```
library(rvest, quietly = TRUE) # uses xml2
```

Attaching package: 'rvest'

The following object is masked from 'package:readr':

```
guess_encoding
```

```

URL <- "https://en.wikipedia.org/wiki/List_of_countries_and_dependencies_by_population"
html <- read_html(URL)
tbls <- html_table(html_elements(html, "table"))
sapply(tbls, nrow)

```

```
[1] 242 12
```

```

pop <- tbls[[1]]
head(pop)

```

```

# A tibble: 6 x 9
  Rank `Country / Dependency` Conti~1 Popul~2 Perce~3 Date Sourc~4 Notes `~`
  <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr> <lg1>
1 - World All 7,973,~ 100% 24 A~ UN pro~ "" NA
2 1 China Asia 1,412,~ 17.7% 31 D~ Offici~ "Exc~ NA
3 2 India Asia 1,360,~ 17.1% 31 D~ Offici~ "The~ NA
4 3 United States North ~ 333,02~ 4.18% 24 A~ Nation~ "The~ NA
5 4 Indonesia Asia[b] 275,77~ 3.46% 1 Ju~ Offici~ "" NA
6 5 Pakistan Asia 235,82~ 2.96% 1 Ju~ UN pro~ "The~ NA
# ... with abbreviated variable names 1: Continent, 2: Population,
# 3: `Percentage of the world`,
# 4: `Source (official or from the United Nations)`

```

(Caution here – notice what format the columns are stored in...)

`read_html()` works by reading in the HTML as text and then parsing it to build up a tree containing the HTML elements. Then `html_elements()` finds the HTML tables and `html_table()` converts them to data frames. `rvest` is part of the tidyverse, so it's often used with piping, e.g.,

```

## Turns out that html_table can take the entire html doc as input
tbls <- URL |> read_html() |> html_table()

```

It's often useful to be able to extract the hyperlinks in an HTML document. We'll find the link using [CSS selectors](#), which allow you to search for elements within HTML:

```

URL <- "http://www1.ncdc.noaa.gov/pub/data/ghcn/daily/by_year"
## approach 1: search for elements with 'href' attribute
links <- read_html(URL) %>% html_elements("[href]") %>% html_attr('href')
## approach 2: search for HTML 'a' tags
links <- read_html(URL) %>% html_elements("a") %>% html_attr('href')
head(links, n = 10)

```

```

[1] "?C=N;O=D"           "?C=M;O=A"           "?C=S;O=A"
[4] "?C=D;O=A"           "/pub/data/ghcn/daily/" "1750.csv.gz"
[7] "1763.csv.gz"         "1764.csv.gz"         "1765.csv.gz"

```

```
[10] "1766.csv.gz"
```

More generally, we may want to read an HTML document, parse it into its components (i.e., the HTML elements), and navigate through the tree structure of the HTML. Here we use the *XPath* language to specify elements rather than CSS selectors. XPath can also be used for navigating through XML documents.

```
## find all 'a' elements that have attribute 'href'; then
## extract the 'href' attribute
links <- read_html(URL) %>% html_elements(xpath = "//a[@href]") %>%
  html_attr('href')
head(links)
```

```
[1] "?C=N;O=D"           "?C=M;O=A"           "?C=S;O=A"
[4] "?C=D;O=A"           "/pub/data/ghcn/daily/" "1750.csv.gz"
```

```
## we can extract various information
listOfANodes <- read_html(URL) %>% html_elements(xpath = "//a[@href]")
listOfANodes %>% html_attr('href') %>% head(n = 10)
```

```
[1] "?C=N;O=D"           "?C=M;O=A"           "?C=S;O=A"
[4] "?C=D;O=A"           "/pub/data/ghcn/daily/" "1750.csv.gz"
[7] "1763.csv.gz"         "1764.csv.gz"         "1765.csv.gz"
[10] "1766.csv.gz"
```

```
listOfANodes %>% html_name() %>% head(n = 10)
```

```
[1] "a" "a" "a" "a" "a" "a" "a" "a" "a" "a"
```

```
listOfANodes %>% html_text() %>% head(n = 10)
```

```
[1] "Name"           "Last modified"   "Size"           "Description"
[5] "Parent Directory" "1750.csv.gz"     "1763.csv.gz"    "1764.csv.gz"
[9] "1765.csv.gz"    "1766.csv.gz"
```

Here's another example of extracting specific components of information from a webpage (results not shown, since headlines will vary from day to day).

```
URL <- "https://www.nytimes.com"
headlines2 <- read_html(URL) %>% html_elements("h2") %>% html_text()
head(headlines2)
headlines3 <- read_html(URL) %>% html_elements("h3") %>% html_text()
head(headlines3)
```

XML

XML is a markup language used to store data in self-describing (no metadata needed) format, often with a hierarchical structure. It consists of sets of elements (also known as nodes because they generally occur in a hierarchical structure and therefore have parents, children, etc.) with tags that identify/name the elements, with some similarity to HTML. Some examples of the use of XML include serving as the underlying format for Microsoft Office and Google Docs documents and for the KML language used for spatial information in Google Earth.

Here's a brief example. The book with id attribute *bk101* is an element; the author of the book is also an element that is a child element of the book. The id attribute allows us to uniquely identify the element.

```
<?xml version="1.0"?>
<catalog>
  <book id="bk101">
    <author>Gambardella, Matthew</author>
    <title>XML Developer's Guide</title>
    <genre>Computer</genre>
    <price>44.95</price>
    <publish_date>2000-10-01</publish_date>
    <description>An in-depth look at creating applications with XML.</description>
  </book>
  <book id="bk102">
    <author>Ralls, Kim</author>
    <title>Midnight Rain</title>
    <genre>Fantasy</genre>
    <price>5.95</price>
    <publish_date>2000-12-16</publish_date>
    <description>A former architect battles corporate zombies, an evil sorceress, and her own ch
  </book>
</catalog>
```

We can read XML documents into R using `xml2::read_xml()` and then manipulate it using other functions from the `xml2` package. Here's an example of working with lending data from the Kiva lending non-profit. You can see the XML format in a browser at

<http://api.kivaws.org/v1/loans/newest.xml>.

XML documents have a tree structure with information at nodes. As above with HTML, one can use the *XPath* language for navigating the tree and finding and extracting information from the node(s) of interest.

Here is some example code for extracting loan info from the Kiva data. We'll first show the 'brute force' approach of working with the data as a list and then the better approach of using XPath.

```
library(xml2)
doc <- read_xml("https://api.kivaws.org/v1/loans/newest.xml")
data <- as_list(doc)
```



```

names(data)

[1] "response"

names(data$response)

[1] "paging" "loans"

length(data$response$loans)

[1] 20

data$response$loans[[2]][c('name', 'activity',
                             'sector', 'location', 'loan_amount')]

$name
$name[[1]]
[1] "Tolgonai"

$activity
$activity[[1]]
[1] "Livestock"

$sector
$sector[[1]]
[1] "Agriculture"

$location
$location$country_code
$location$country_code[[1]]
[1] "KG"

$location$country
$location$country[[1]]
[1] "Kyrgyzstan"

$location$town
$location$town[[1]]
[1] "Aksy district, Jalalabad region"

```

```
$location$geo
$location$geo$level
$location$geo$level[[1]]
[1] "town"
```

```
$location$geo$pairs
$location$geo$pairs[[1]]
[1] "41.673257 72.046816"
```

```
$location$geo$type
$location$geo$type[[1]]
[1] "point"
```

```
$loan_amount
$loan_amount[[1]]
[1] "3100"
```

```
## alternatively, extract only the 'loans' info (and use pipes)
loansNode <- doc %>% html_elements('loans')
loanInfo <- loansNode %>% xml_children() %>% as_list()
length(loanInfo)
```

```
[1] 20
```

```
names(loanInfo[[1]])
```

```
[1] "id" "name"
[3] "description" "status"
[5] "funded_amount" "basket_amount"
[7] "image" "activity"
[9] "sector" "themes"
[11] "use" "location"
[13] "partner_id" "posted_date"
[15] "planned_expiration_date" "loan_amount"
[17] "borrower_count" "lender_count"
[19] "bonus_credit_eligibility" "tags"
```

```
names(loanInfo[[1]]$location)
```

```

[1] "country_code" "country"      "town"      "geo"

## suppose we only want the country locations of the loans (using XPath)
xml_find_all(loansNode, '//location//country')

{xml_nodeset (20)}
[1] <country>Lesotho</country>
[2] <country>Kyrgyzstan</country>
[3] <country>Lesotho</country>
[4] <country>Lesotho</country>
[5] <country>Lesotho</country>
[6] <country>Lesotho</country>
[7] <country>Lesotho</country>
[8] <country>Lesotho</country>
[9] <country>Philippines</country>
[10] <country>Philippines</country>
[11] <country>Tajikistan</country>
[12] <country>Philippines</country>
[13] <country>Philippines</country>
[14] <country>Philippines</country>
[15] <country>Philippines</country>
[16] <country>Philippines</country>
[17] <country>Burkina Faso</country>
[18] <country>Burkina Faso</country>
[19] <country>Lesotho</country>
[20] <country>Uganda</country>

xml_find_all(loansNode, '//location//country') %>% xml_text()

[1] "Lesotho"      "Kyrgyzstan"   "Lesotho"      "Lesotho"      "Lesotho"
[6] "Lesotho"      "Lesotho"      "Lesotho"      "Philippines"  "Philippines"
[11] "Tajikistan"   "Philippines"  "Philippines"  "Philippines"  "Philippines"
[16] "Philippines"  "Burkina Faso" "Burkina Faso" "Lesotho"      "Uganda"

## or extract the geographic coordinates
xml_find_all(loansNode, '//location//geo/pairs')

{xml_nodeset (20)}
[1] <pairs>-29.315077 27.486923</pairs>
[2] <pairs>41.673257 72.046816</pairs>
[3] <pairs>-29.315077 27.486923</pairs>
[4] <pairs>-29.681724 27.402676</pairs>
[5] <pairs>-29.601652 27.120298</pairs>
[6] <pairs>-29.523898 28.608975</pairs>
[7] <pairs>-29.400969 27.763725</pairs>

```

```

[8] <pairs>-29.523898 28.608975</pairs>
[9] <pairs>10.432923 124.994751</pairs>
[10] <pairs>10.432923 124.994751</pairs>
[11] <pairs>38.598056 68.977778</pairs>
[12] <pairs>10.432923 124.994751</pairs>
[13] <pairs>10.432923 124.994751</pairs>
[14] <pairs>10.432923 124.994751</pairs>
[15] <pairs>12.619862 125.016787</pairs>
[16] <pairs>12.619862 125.016787</pairs>
[17] <pairs>12.371428 -1.51966</pairs>
[18] <pairs>12.371428 -1.51966</pairs>
[19] <pairs>-29.779915 27.288186</pairs>
[20] <pairs>1.064444 34.179444</pairs>

```

JSON

JSON files are structured as “attribute-value” pairs (aka “key-value” pairs), often with a hierarchical structure. Here’s a brief example:

```

{
  "firstName": "John",
  "lastName": "Smith",
  "isAlive": true,
  "age": 25,
  "address": {
    "streetAddress": "21 2nd Street",
    "city": "New York",
    "state": "NY",
    "postalCode": "10021-3100"
  },
  "phoneNumbers": [
    {
      "type": "home",
      "number": "212 555-1234"
    },
    {
      "type": "office",
      "number": "646 555-4567"
    }
  ],
  "children": [],
  "spouse": null
}

```

A set of key-value pairs is a named array and is placed inside braces (squiggly brackets). Note the nestedness of arrays within arrays (e.g., address within the overarching person array and the use of square brackets for unnamed arrays (i.e., vectors of information), as well as the use of different types:

character strings, numbers, null, and (not shown) boolean/logical values. JSON and XML can be used in similar ways, but JSON is less *verbose* than XML.

We can read JSON into R using *fromJSON()* in the *jsonlite* package. Let's play again with the Kiva data. The same data that we had worked with in XML format is also available in JSON format: <http://api.kivaws.org/v1/loans/newest.json>.

```
library(jsonlite)
data <- fromJSON("http://api.kivaws.org/v1/loans/newest.json")
class(data)
```

```
[1] "list"
```

```
names(data)
```

```
[1] "paging" "loans"
```

```
class(data$loans) # nice!
```

```
[1] "data.frame"
```

```
head(data$loans)
```

	id	name	languages	status	funded_amount	basket_amount
1	2430686	Motsukunyane	en	fundraising	0	0
2	2432992	Tolgonai	ru, en	fundraising	0	0
3	2430680	Khotso	en	fundraising	0	25
4	2430652	Matebello	en	fundraising	25	0
5	2430659	Makeletso	en	fundraising	25	0
6	2430662	Masehlomeng	en	fundraising	0	0

	image.id	image.template_id	activity	sector
1	4925389	1	Home Energy	Personal Use
2	4928728	1	Livestock	Agriculture
3	4925368	1	Home Energy	Personal Use
4	4925270	1	Home Energy	Personal Use
5	4925286	1	Home Energy	Personal Use
6	4925296	1	Home Energy	Personal Use

	themes
1	Green, Earth Day Campaign, Clean Energy
2	Green, Earth Day Campaign
3	Green, Earth Day Campaign, Clean Energy
4	Green, Earth Day Campaign, Clean Energy
5	Green, Earth Day Campaign, Clean Energy
6	Green, Earth Day Campaign, Clean Energy

```

1 to purchase the ACE One hybrid energy system for cooking, lighting and charging in a way that is sa
2 to purchase livestock to increase her in
3 to purchase the ACE One hybrid energy system for cooking, lighting and charging in a way that is sa
4 to purchase the ACE One hybrid energy system for cooking, lighting and charging in a way that is sa
5 to purchase the ACE One hybrid energy system for cooking, lighting and charging in a way that is sa
6 to purchase the ACE One hybrid energy system for cooking, lighting and charging in a way that is sa
  location.country_code location.country location.town
1 LS Lesotho Popa Tholang, Maseru
2 KG Kyrgyzstan Aksy district, Jalalabad region
3 LS Lesotho Popa Tholang, Maseru
4 LS Lesotho Mahareng, Mafeteng
5 LS Lesotho Thabana mohlomi, Mafeteng
6 LS Lesotho Methalaneng, Thaba-Tseka
  location.geo.level location.geo.pairs location.geo.type partner_id
1 town -29.315077 27.486923 point 452
2 town 41.673257 72.046816 point 171
3 town -29.315077 27.486923 point 452
4 town -29.681724 27.402676 point 452
5 town -29.601652 27.120298 point 452
6 town -29.523898 28.608975 point 452
  posted_date planned_expiration_date loan_amount borrower_count
1 2022-08-26T18:50:08Z 2022-09-30T18:50:08Z 100 1
2 2022-08-26T18:50:03Z 2022-09-30T18:50:03Z 3100 1
3 2022-08-26T18:30:12Z 2022-09-30T18:30:11Z 100 1
4 2022-08-26T18:30:11Z 2022-09-30T18:30:11Z 100 1
5 2022-08-26T18:30:11Z 2022-09-30T18:30:11Z 100 1
6 2022-08-26T18:30:11Z 2022-09-30T18:30:11Z 100 1
  lender_count bonus_credit_eligibility
1 0 FALSE
2 0 TRUE
3 0 FALSE
4 1 FALSE
5 1 FALSE
6 0 FALSE

1 volunteer_pick, #Eco-friendly, #Health and Sanitation, #Technology, #Repair Renew Repl
2 #Sustainable Ag, #Animals, #Parent, #Repeat Bo
3 volunteer_pick, #Eco-friendly, #Health and Sanitation, #Technology, #Repair Renew Repl
4 volunteer_pick, #Eco-friendly, #Health and Sanitation, #Technology, #Repair Renew Repl
5 volunteer_pick, #Eco-friendly, #Parent, #Health and Sanitation, #Technology, #Repair Renew Replace,
6 volunteer_pick, #Eco-friendly, #Health and Sanitation, #Technology, #Repair Renew Repl

```

```
data$loans[1, 'location.geo.pairs'] # hmmm...
```

NULL

```
data$loans[1, 'location']
```

```
country_code country          town geo.level      geo.pairs
1          LS Lesotho Popa Tholang, Maseru      town -29.315077 27.486923
geo.type
1    point
```

One disadvantage of JSON is that it is not set up to deal with missing values, infinity, etc.

Webscraping and web APIs

Here we'll see some examples of making requests over the Web to get data. We'll use APIs to systematically query a website for information. Ideally, but not always, the API will be documented. In many cases that simply amounts to making an HTTP GET request, which is done by constructing a URL.

The packages *RCurl* and *httr* are useful for a wide variety of such functionality. Note that much of the functionality I describe below is also possible within the shell using either *wget* or *curl*.

Webscraping ethics and best practices

Webscraping is the process of extracting data from the web, either directly from a website or using a web API (application programming interface).

1. **Should you webscrape?** In general, if we can avoid webscraping (particularly if there is not an API) and instead directly download a data file from a website, that is greatly preferred.
2. **May you webscrape?** Before you set up any automated downloading of materials/data from the web you should make sure that what you are about to do is consistent with the rules provided by the website.

Some places to look for information on what the website allows are:

- legal pages such as Terms of Service or Terms and Conditions on the website.
- check the robots.txt file (e.g., <https://scholar.google.com/robots.txt>) to see what a web crawler is allowed to do, and whether the site requires a particular delay between requests to the sites
- potentially contact the site owner if you plan to scrape a large amount of data

Here are some links with useful information:

- [A blog post overview on webscraping and robots.txt](#)
- [Blog post on webscraping ethics](#)
- [Some information on how to understand a robots.txt file](#)

Tips for when you make automated requests:

- When debugging code that processes the result of such a request, just run the request once, save (i.e., cache) the result, and then work on the processing code applied to the result. Don't make the same request over and over again.

- In many cases you will want to include a time delay between your automated requests to a site, including if you are not actually crawling a site but just want to automate a small number of queries.

What is HTTP?

HTTP (hypertext transfer protocol) is a system for communicating information from a server (i.e., the website of interest) to a client (e.g., your laptop). The client sends a request and the server sends a response.

When you go to a website in a browser, your browser makes an HTTP GET request to the website. Similarly, when we did some downloading of html from webpages above, we used an HTTP GET request.

Anytime the URL you enter includes parameter information after a question mark (`www.somewebsite.com?param1=arg1&p` you are using an API.

The response to an HTTP request will include a status code, which can be interpreted based on [this information](#).

The response will generally contain content in the form of text (e.g., HTML, XML, JSON) or raw bytes.

APIs: REST- and SOAP-based web services

Ideally a web service documents their API (Applications Programming Interface) that serves data or allows other interactions. REST and SOAP are popular API standards/styles. Both REST and SOAP use HTTP requests; we'll focus on REST as it is more common and simpler. When using REST, we access *resources*, which might be a Facebook account or a database of stock quotes. The API will (hopefully) document what information it expects from the user and will return the result in a standard format (often a particular file format rather than producing a webpage).

Often the format of the request is a URL (aka an endpoint) plus a query string, passed as a GET request. Let's search for plumbers near Berkeley, and we'll see the GET request, in the form:

https://www.yelp.com/search?find_desc=plumbers&find_loc=Berkeley+CA&ns=1

- the query string begins with ?
- there are one or more **Parameter=Argument** pairs
- pairs are separated by &
- + is used in place of each space

Let's see an example of accessing economic data from the World Bank, using the [documentation for their API](#). Following the [API call structure](#), we can download (for example), data on various countries. The documentation indicates that our REST-based query can use either a URL structure or an argument-based structure.


```

## Queries based on the documentation
api_url <- "http://api.worldbank.org/V2/incomeLevel/LIC/country"
api_args <- "http://api.worldbank.org/V2/country?incomeLevel=LIC"

## Generalizing a bit
req <- "http://api.worldbank.org/V2/country?incomeLevel=MIC&format=json"
data <- fromJSON(req)
## Be careful of data truncation/pagination
req <- "http://api.worldbank.org/V2/country?incomeLevel=MIC&format=json&per_page=1000"
data <- fromJSON(req)

## Programmatic control
baseURL <- "http://api.worldbank.org/V2/country"
group <- 'MIC'
format <- 'json'
args <- c(incomeLevel = group, format = format, per_page = 1000)
url <- paste0(baseURL, "?",
  paste( paste(names(args), args, sep = "="), collapse = "&"))

data <- fromJSON(url)
class(data)

```

```
[1] "list"
```

```
length(data)
```

```
[1] 2
```

```
names(data)
```

```
NULL
```

```
head(data[[2]])
```

	id	iso2Code	name	region.id	region.iso2code
1	AGO	AO	Angola	SSF	ZG
2	ALB	AL	Albania	ECS	Z7
3	ARG	AR	Argentina	LCN	ZJ
4	ARM	AM	Armenia	ECS	Z7
5	ASM	AS	American Samoa	EAS	Z4
6	AZE	AZ	Azerbaijan	ECS	Z7

	region.value	adminregion.id	adminregion.iso2code
1	Sub-Saharan Africa	SSA	ZF
2	Europe & Central Asia	ECA	7E
3	Latin America & Caribbean	LAC	XJ

4	Europe & Central Asia	ECA	7E
5	East Asia & Pacific	EAP	4E
6	Europe & Central Asia	ECA	7E

	adminregion.value	incomeLevel.id
1	Sub-Saharan Africa (excluding high income)	LMC
2	Europe & Central Asia (excluding high income)	UMC
3	Latin America & Caribbean (excluding high income)	UMC
4	Europe & Central Asia (excluding high income)	UMC
5	East Asia & Pacific (excluding high income)	UMC
6	Europe & Central Asia (excluding high income)	UMC

	incomeLevel.iso2code	incomeLevel.value	lendingType.id	lendingType.iso2code
1	XN	Lower middle income	IBD	XF
2	XT	Upper middle income	IBD	XF
3	XT	Upper middle income	IBD	XF
4	XT	Upper middle income	IBD	XF
5	XT	Upper middle income	LNK	XX
6	XT	Upper middle income	IBD	XF

	lendingType.value	capitalCity	longitude	latitude
1	IBRD	Luanda	13.242	-8.81155
2	IBRD	Tirane	19.8172	41.3317
3	IBRD	Buenos Aires	-58.4173	-34.6118
4	IBRD	Yerevan	44.509	40.1596
5	Not classified	Pago Pago	-170.691	-14.2846
6	IBRD	Baku	49.8932	40.3834

APIs can change and disappear. Last year the example above involved the World Bank's Climate Data API, which I can no longer find!

As another example, here we can see the [US Treasury Department API](#), which allows us to construct queries for federal financial data.

The Nolan and Temple Lang book provides a number of examples of different ways of authenticating with web services that control access to the service.

Finally, some web services allow us to pass information to the service in addition to just getting data or information. E.g., you can programmatically interact with your Facebook, Dropbox, and Google Drive accounts using REST based on HTTP POST, PUT, and DELETE requests. Authentication is of course important in these contexts and some times you would first authenticate with your login and password and receive a "token". This token would then be used in subsequent interactions in the same session.

I created your *github.berkeley.edu* accounts from Python by interacting with the [GitHub API](#) using the *requests* package.

HTTP requests by deconstructing an (undocumented) API

In some cases an API may not be documented or we might be lazy and not use the documentation. Instead we might deconstruct the queries a browser makes and then mimic that behavior, in some

cases having to parse HTML output to get at data. Note that if the webpage changes even a little bit, our carefully constructed query syntax may fail.

Let's look at some UN data (agricultural crop data). By going to <http://data.un.org/Explorer.aspx?d=FAO>, and clicking on "Crops", we'll see a bunch of agricultural products with "View data" links. Click on "apricots" as an example and you'll see a "Download" button that allows you to download a CSV of the data. Let's select a range of years and then try to download "by hand". Sometimes we can right-click on the link that will download the data and directly see the URL that is being accessed and then one can deconstruct it so that you can create URLs programmatically to download the data you want.

In this case, we can't see the full URL that is being used because there's some Javascript involved. Therefore, rather than looking at the URL associated with a link we need to view the actual HTTP request sent by our browser to the server. We can do this using features of the browser (e.g., in Firefox see **Web Developer** -> **Network** and in Chrome **View** -> **Developer** -> **Developer tools** and choose the **Network** tab) (or right-click on the webpage and select **Inspect** and then **Network**). Based on this we can see that an HTTP GET request is being used with a URL such as:

<http://data.un.org/Handlers/DownloadHandler.ashx?DataFilter=itemCode:526;year:2012,2013,2014,2015,2016,2017&DataMartId=FAO&Format=csv&c=2,4,5,6,7&s=countryName:asc,elementCode:asc,year:desc>.

We're now able to easily download the data using that URL, which we can fairly easily construct using string processing in bash, R, or Python, such as this (here I just paste it together directly, but using more structured syntax such as I used for the World Bank example would be better):

```
## example URL:
## http://data.un.org/Handlers/DownloadHandler.ashx?DataFilter=itemCode:526;
##year:2012,2013,2014,2015,2016,2017&DataMartId=FAO&Format=csv&c=2,4,5,6,7&
##s=countryName:asc,elementCode:asc,year:desc
itemCode <- 526
baseURL <- "http://data.un.org/Handlers/DownloadHandler.ashx"
yrs <- paste(as.character(2012:2017), collapse = ",")
filter <- paste0("?DataFilter=itemCode:", itemCode, ";year:", yrs)
args1 <- "&DataMartId=FAO&Format=csv&c=2,3,4,5,6,7&"
args2 <- "s=countryName:asc,elementCode:asc,year:desc"
url <- paste0(baseURL, filter, args1, args2)
## if the website provided a CSV we could just do this:
## apricots <- read.csv(url)
## but it zips the file
temp <- tempfile() ## give name for a temporary file
download.file(url, temp)
dat <- read.csv(unzip(temp)) ## using a connection (see Section 2)

head(dat)
```

	Country.or.Area	Element.Code	Element
1	Afghanistan	432	Gross Production Index Number (2014–2016 = 100)

```

2      Afghanistan      432 Gross Production Index Number (2014-2016 = 100)
3      Afghanistan      432 Gross Production Index Number (2014-2016 = 100)
4      Afghanistan      432 Gross Production Index Number (2014-2016 = 100)
5      Afghanistan      432 Gross Production Index Number (2014-2016 = 100)
6      Afghanistan      432 Gross Production Index Number (2014-2016 = 100)
  Year Unit Value Value.Footnotes
1 2017 index 202.19              Fc
2 2016 index  27.45              Fc
3 2015 index 134.50              Fc
4 2014 index 138.05              Fc
5 2013 index 138.05              Fc
6 2012 index 128.08              Fc

```

So, what have we achieved?

1. We have a reproducible workflow we can share with others (perhaps ourself in the future).
2. We can automate the process of downloading many such files.

More details on HTTP requests

A more sophisticated way to do the download is to pass the request in a structured way with named input parameters. This request is easier to construct programmatically. Here what is returned is a zip file, which is represented in R as a sequence of “raw” bytes. We can use `httr`’s `GET()`, followed by writing to disk and reading back in, as follows (for some reason knitr won’t print the output...):

```
library(httr)
```

Attaching package: ‘httr’

The following object is masked from ‘package:curl’:

```
handle_reset
```

```

output2 <- GET(baseUrl, query = list(
  DataFilter = paste0("itemCode:", itemCode, ";year:", yrs),
  DataMartID = "FAO", Format = "csv", c = "2,3,4,5,6,7",
  s = "countryName:asc,elementCode:asc,year:desc"))
temp <- tempfile() ## give name for a temporary file
writeBin(content(output2, 'raw'), temp) ## write out as zip file
dat <- read.csv(unzip(temp))
head(dat)

```

	Country.or.Area	Element.Code	Element
1	Afghanistan	432 Gross Production Index Number (2014-2016 = 100)	
2	Afghanistan	432 Gross Production Index Number (2014-2016 = 100)	
3	Afghanistan	432 Gross Production Index Number (2014-2016 = 100)	
4	Afghanistan	432 Gross Production Index Number (2014-2016 = 100)	

5	Afghanistan		432	Gross Production Index Number (2014-2016 = 100)
6	Afghanistan		432	Gross Production Index Number (2014-2016 = 100)
	Year	Unit	Value	Value.Footnotes
1	2017	index	202.19	Fc
2	2016	index	27.45	Fc
3	2015	index	134.50	Fc
4	2014	index	138.05	Fc
5	2013	index	138.05	Fc
6	2012	index	128.08	Fc

In some cases we may need to send a lot of information as part of the URL in a GET request. If it gets to be too long (e.g., more than 2048 characters) many web servers will reject the request. Instead we may need to use an HTTP POST request (POST requests are often used for submitting web forms). A typical request would have syntax like this search (using *RCurl*):

```
if(url.exists('http://www.wormbase.org/db/searches/advanced/dumper')) {
  x = postForm('http://www.wormbase.org/db/searches/advanced/dumper',
    species="briggsae",
    list="",
    flank3="0",
    flank5="0",
    feature="Gene Models",
    dump = "Plain TEXT",
    orientation = "Relative to feature",
    relative = "Chromosome",
    DNA ="flanking sequences only",
    .cgifields = paste(c("feature", "orientation", "DNA",
      "dump","relative"), collapse=", "))
}
```

Unfortunately that specific search doesn't work because the server URL and/or API seem to have changed. But it gives you an idea of what the format would look like.

httr and *RCurl* can handle other kinds of HTTP requests such as PUT and DELETE. Finally, some websites use cookies to keep track of users and you may need to download a cookie in the first interaction with the HTTP server and then send that cookie with later interactions. More details are available in the Nolan and Temple Lang book.

Packaged access to an API

For popular websites/data sources, a developer may have packaged up the API calls in a user-friendly fashion for use from R, Python or other software. For example there are Python (*twitter*) and R (*twitterR*) packages for interfacing with Twitter via its API.

Here's some example code for Python (the Python package seems to be more fully-featured than the R package). This looks up the US senators' Twitter names and then downloads a portion of each of their timelines, i.e., the time series of their tweets. Note that Twitter has limits on how much one can download at once.

```

import json
import twitter

# You will need to set the following variables with your
# personal information. To do this you will need to create
# a personal account on Twitter (if you don't already have
# one). Once you've created an account, create a new
# application here:
#   https://dev.twitter.com/apps
#
# You can manage your applications here:
#   https://apps.twitter.com/
#
# Select your application and then under the section labeled
# "Key and Access Tokens", you will find the information needed
# below. Keep this information private.
CONSUMER_KEY      = ""
CONSUMER_SECRET   = ""
OAUTH_TOKEN       = ""
OAUTH_TOKEN_SECRET = ""

auth = twitter.oauth.OAuth(OAUTH_TOKEN, OAUTH_TOKEN_SECRET,
                           CONSUMER_KEY, CONSUMER_SECRET)
api = twitter.Twitter(auth=auth)

# get the list of senators
senators = api.lists.members(owner_screen_name="gov", slug="us-senate", count=100)

# get all the senators' timelines
names = [d["screen_name"] for d in senators["users"]]
timelines = [api.statuses.user_timeline(screen_name=name, count = 500)
              for name in names]

# save information out to JSON
with open("senators-list.json", "w") as f:
    json.dump(senators, f, indent=4, sort_keys=True)
with open("timelines.json", "w") as f:
    json.dump(timelines, f, indent=4, sort_keys=True)

```

Accessing dynamic pages

Some websites dynamically change in reaction to the user behavior. In these cases you need a tool that can mimic the behavior of a human interacting with a site. Some options are:

- *selenium* (and the *RSelenium* wrapper for R) is a popular tool for doing this.

- *splash* (and the *splashr* wrapper for R) is another approach.
- *htmlunit* is another tool for this.

5. File and string encodings

Text (either in the form of a file with regular language in it or a data file with fields of character strings) will often contain characters that are not part of the [limited ASCII set of characters](#), which has $2^7 = 128$ characters and control codes; basically what you see on a standard US keyboard. Each character takes up one byte (8 bits) of space (there is an unused bit that comes in handy in the UTF-8 context). We can actually hand-generate an ASCII file using the binary representation of each character in R as an illustration.

The letter “M” is encoded based on the ASCII standard in bits as “01001101” as seen in the link above. For convenience, this is often written as two base-16 numbers (i.e., hexadecimal), where “0100”=“4” and “1101”=“d”, hence we have “4d” in hexadecimal.

```
## 4d in hexadecimal is 'M'
## 0a is a newline (at least in Linux/Mac)
## "0x" is how we tell R we are using hexadecimal
x <- as.raw(c('0x4d', '0x6f', '0x6d', '0x0a')) ## i.e., "Mom\n" in ascii
x
```

```
[1] 4d 6f 6d 0a
```

```
charToRaw('Mom\n')
```

```
[1] 4d 6f 6d 0a
```

```
writeBin(x, 'tmp.txt')
readLines('tmp.txt')
```

```
[1] "Mom"
```

```
system('ls -l tmp.txt', intern = TRUE)
```

```
[1] "-rw-r--r-- 1 paciorek scfstaff 4 Aug 26 11:54 tmp.txt"
```

```
system('cat tmp.txt')
```

When encountering non-ASCII files, in some cases you may need to deal with the text encoding (the mapping of individual characters (including tabs, returns, etc.) to a set of numeric codes). There are a variety of different encodings for text files, with different ones common on different operating systems. [UTF-8](#) is an encoding for the Unicode characters that includes more than 110,000 characters from 100 different alphabets/scripts. It’s widely used on the web. Latin-1 encodes a small subset of Unicode

and contains the characters used in many European languages (e.g., letters with accents). Here's an example of using a non-ASCII Unicode character:

```
## n-tilde and division symbol as Unicode 'code points'
x2_unicode <- 'Pe\u00f1a 3\u00f72'
Encoding(x2_unicode)

[1] "UTF-8"

x2_unicode

[1] "Peña 3÷2"

charToRaw(x2_unicode)

[1] 50 65 c3 b1 61 20 33 c3 b7 32

charToRaw('\u00f1')      # indeed - two bytes, not one

[1] c3 b1

## specified directly as hexadecimal in UTF-8 encoding
x2_utf8 <- 'Pe\xc3\xb1a 3\xc3\xb72'
x2_utf8

[1] "Peña 3÷2"

writeBin(x2_unicode, 'tmp2.txt')
## Here n-tilde and division symbol take up two bytes
## but there is an extraneous null byte in there; not sure why.
system('ls -l tmp2.txt')
## The system knows how to interpret the UTF-8 encoded file
## and represent the Unicode character on the screen:
system('cat tmp2.txt')
```

UTF-8 is cleverly designed in terms of the bit-wise representation of characters such that ASCII characters still take up one byte, and most other characters take two bytes, but some take four bytes. In fact it is even more clever than that - the representation is such that the bits of a one-byte character never appear within the representation of a two- or three- or four-byte character (and similarly for two-byte characters in three- or four-byte characters, etc.).

The UNIX utility *file*, e.g. `file tmp.txt` can help provide some information. `read.table()` in R takes arguments *fileEncoding* and *encoding* that allow one to specify the encoding as one reads text in. The UNIX utility *iconv* and the R function `iconv()` can help with conversions.

In US installations of R, the default encoding is UTF-8; note below that various types of information are interpreted in US English with the encoding UTF-8:

```
Sys.getlocale()
```

```
[1] "LC_CTYPE=en_US.UTF-8;LC_NUMERIC=C;LC_TIME=en_US.UTF-8;LC_COLLATE=en_US.UTF-8;LC_MONETARY=en_US.U
```

With strings already in R, you can convert between encodings with *iconv()*:

```
text <- "Melhore sua seguran\xe7a"  
Encoding(text)
```

```
[1] "unknown"
```

```
Encoding(text) <- "latin1"  
text ## this prints out correctly in R, but is not correct in the PDF
```

```
[1] "Melhore sua segurança"
```

```
text <- "Melhore sua seguran\xe7a"  
textUTF8 <- iconv(text, from = "latin1", to = "UTF-8")  
Encoding(textUTF8)
```

```
[1] "UTF-8"
```

```
textUTF8
```

```
[1] "Melhore sua segurança"
```

```
iconv(text, from = "latin1", to = "ASCII", sub = "???")
```

```
[1] "Melhore sua seguran???a"
```

You can also mark a string with an encoding, so R knows how to display it correctly (again, this prints out incorrectly in the PDF):

```
x <- "fa\xe7ile"  
Encoding(x) <- "latin1"  
x
```

```
[1] "façile"
```

```
## playing around...
x <- "\xa1\xa2\xa3\xfa\xfb"
Encoding(x) <- "latin1"
x
```

```
[1] "ï ¢ £ ñ ò"
```

An R error message with multi-byte string in the message often indicates an encoding issue. In particular errors often arise when trying to do string manipulations in R on character vectors for which the encoding is not properly set. Here's an example with some Internet logging data that we used a few years ago in class in a problem set and which caused some problems.

```
load('../data/IPs.RData') # loads in an object named 'text'
tmp <- try(substring(text, 1, 15))
```

```
Error in substring(text, 1, 15) :
  invalid multibyte string at '<bf>a7lw8<6c>z2nX,%@ [128.32.244.179] by ncp-email with ESMT
(SMTDP32-7.04) id A06E24A0116; Mon, 10 Jun 2002 11:43:42 +0800'
```

```
## the issue occurs with the 6402th element (found by trial and error):
tmp <- try(substring(text[1:6401], 1, 15))
tmp <- try(substring(text[1:6402], 1, 15))
```

```
Error in substring(text[1:6402], 1, 15) :
  invalid multibyte string at '<bf>a7lw8<6c>z2nX,%@ [128.32.244.179] by ncp-email with ESMT
(SMTDP32-7.04) id A06E24A0116; Mon, 10 Jun 2002 11:43:42 +0800'
```

```
text[6402] # note the Latin-1 character
```

```
[1] "from 5#c\xbfa7lw8lz2nX,%@ [128.32.244.179] by ncp-email with ESMT\n(SMTDP32-7.04) id A06E24A0116; Mon, 10 Jun 2002 11:43:42 +0800"
```

```
table(Encoding(text))
```

```
unknown
6936
```

```
## Option 1
Encoding(text) <- "latin1"
tmp <- try(substring(text, 1, 15))
tmp[6402]
```

```
[1] "from 5#c;a7lw8l"
```

```
## Option 2
load('../data/IPs.RData') # loads in an object named 'text'
tmp <- try(substring(text, 1, 15))
```

```
Error in substring(text, 1, 15) :
  invalid multibyte string at '<bf>a7lw8<6c>z2nX,%@ [128.32.244.179] by ncpc-email with ESMTP
(SMTPD32-7.04) id A06E24A0116; Mon, 10 Jun 2002 11:43:42 +0800'
```

```
text <- iconv(text, from = "latin1", to = "UTF-8")
tmp <- try(substring(text, 1, 15))
```

6. Data structures

As we're reading data into R or other languages, it's important to think about the data structures we'll use to store the information. The data structure we choose can affect:

- the amount of memory we need,
- how quickly we can access the information in the data structure,
- how much copying needs to be done to add information to or remove information from the data structure,
- how efficiently we can use the data in subsequent computations.

This means that what you plan to do with the data should guide what kind of structure you store the data in.

Standard data structures in R and Python

- In R and Python, one often ends up working with dataframes, lists, and vectors/matrices/arrays/tensors.
- In R, if we are not working with rectangular datasets or standard numerical objects, we often end up using lists or enhanced versions of lists, sometimes with deeply nested structures.
- In Python we commonly work with data structures that are part of additional packages, in particular numpy arrays and pandas dataframes.
- Dictionaries in Python allow for easy use of key-value pairs where one can access values based on their key/label. In R one can do something similar with named vectors or named lists or (more efficiently) by using environments.

In Unit 7, we'll talk about *distributed* data structures that allow one to easily work with data distributed across multiple computers.

Other kinds of data structures

You may have heard of various other kinds of data structures, such as linked lists, trees, graphs, queues, and stacks. One of the key aspects that differentiate such data structures is how one navigates through the elements.

Sets are collections of elements that don't have any duplicates (like a mathematical set).

With a *linked list*, with each element (or node) has a value and a pointer (reference) to the location of the next element. (With a doubly-linked list, there is also a pointer back to the previous element.) One big advantage of this is that one can insert an element by simply modifying the pointers involved at the site of the insertion, without copying any of the other elements in the list. A big disadvantage is that to get to an element you have to navigate through the list.




Figure 1: Linked list (courtesy of computersciencewiki.org)

Both *trees* and *graphs* are collections of nodes (vertices) and links (edges). A tree involves a set of nodes and links to child nodes (also possibly containing information linking the child nodes to their parent nodes). With a graph, the links might not be directional, and there can be cycles.




Figure 2: Tree (courtesy of computersciencewiki.org)

A *stack* is a collection of elements that behave like a stack of lunch trays. You can only access the



Figure 3: Graph (courtesy of [computersciencewiki.org](https://en.wikipedia.org/wiki/Computer_Science_Wiki))

top element directly (“last in, first out”), so the operations are that you can push a new element onto the stack or pop the top element off the stack. In fact, nested function calls behave as stacks, and the memory used in the process of evaluating the function calls is called the ‘stack’.

A *queue* is like the line at a grocery store, behaving as “first in, first out”.

One can use such data structures either directly or via add-on packages in R and Python, though I don’t think they’re all that commonly used in R. This is probably because statistical/data science/machine learning workflows often involve either ‘rectangular’ data (i.e., dataframe-style data) and/or mathematical computations with arrays. That said, trees and graphs are widely used.

Some related concepts that we’ll discuss further in Unit 5 include:

- types: this refers to how a given piece of information is stored and what operations can be done with the information.
 - ‘primitive’ types are the most basic types that often relate directly to how data are stored in memory or on disk (e.g., booleans, integers, numeric (real-valued), character, pointer (address, reference)).
- pointers: references to other locations (addresses) in memory. One often uses pointers to avoid unnecessary copying of data.
- hashes: hashing involves fast lookup of the value associated with a key (a label), using a hash function, which allows one to convert the key to an address. This avoids having to find the value associated with a specific key by looking through all the keys until the key of interest is found (an $O(n)$ operation).